







Methods of Artificial Intelligence

Local Search


A background image showing a human hand holding a robotic arm, symbolizing the transition from classical to local search. The hand is on the left, and the robotic arm is on the right, with the text centered over the intersection. A thick red curved line is at the top of the slide.

From Classical Search to Local Search


Search Problems



Search Space and Search Tree



search space



Classical Search Paradigms

□ Uninformed Search

- No information about search space topology (blind search)

□ Informed Search

- We can estimate distance to goal states

Classical Search Algorithms

☐ Uninformed (blind)

■ Random search

■ Systematic search

- ☐ Depth-first search (DFS)
- ☐ Breadth-first search (BFS)
- ☐ Variations
 - Uniform-cost search
 - Depth-limited search

■ Features

- ☐ no information about the path cost from the current state to the goal state

☐ Informed (heuristic)

■ Systematic search

- ☐ Greedy best-first search
- ☐ A* search

■ Features



- ☐ heuristic information about the path cost from current state to goal state (background knowledge)

From Classical Search to Local Search

□ In **classical search**, we look for an (optimal) sequence of actions that leads from an **initial state** to a **goal state**

□ However, in many problems, there is

- no given initial state (CSPs)
- no obvious transition model
- no simple goal test (optimization problems)



□ In particular, path from initial to goal state is often irrelevant

□ **Local Search** is often better suited to address these problems

Local Search Main Ideas

- States often correspond to **variable assignments**
- We define **neighborhood** that determines what other assignments can be reached from a given assignment
- Local Search works roughly as follows:
 - start from some **initial assignment** (e.g. random choice)
 - select a promising **neighbor** of the current assignment
 - continue selecting neighbors until we cannot improve anymore

Local Search Applications

- (Discrete) Optimization Problems
 - Production planning (maximize profit)
 - Scheduling (minimize conflicts)
 - Machine Learning (minimize classification error)
 - ...
- Constraint Satisfaction Problems

Algorithms Overview

We will discuss the following local search algorithms:


- ☐ Hill-Climbing
- ☐ Simulated Annealing
- ☐ Local Beam Search
- ☐ Genetic Algorithms

A hand holding a game controller is visible in the background, slightly faded. A thick red curved line arches across the top of the slide. The text is centered in a dark red serif font.


State-space Landscape and Hill-Climbing

State-space Landscape

- Suppose, we want to maximize real function $f(x)$ (**objective function**)
- we can do so by starting from random point and following ascent direction (**gradient ascent algorithm**)




Challenges



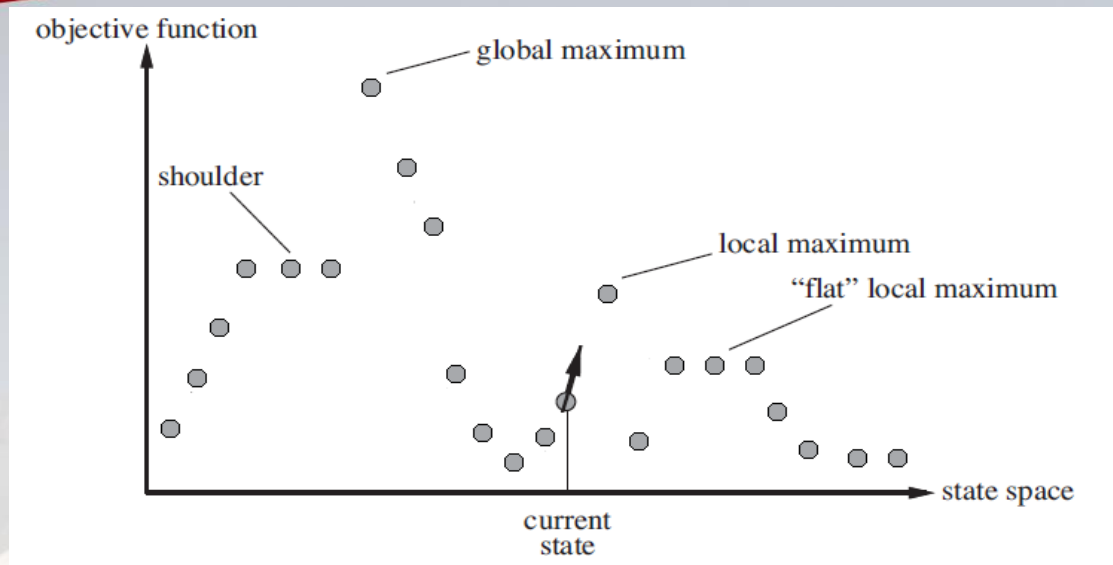
- if we stop if no ascent is possible anymore, we may end up in a non-global local maximum or plateaux
- If we try to reach end of plateaux, we may search forever

Discrete State Spaces



- We will often consider discrete state spaces here
(we cannot move continuously through state space)
- Often our state space is even finite, but exponentially large
- However, some problems of continuous spaces remain

Navigating in Discrete State Spaces



- In (unconstrained) continuous state spaces, we can move in an arbitrary direction for an arbitrary length
- In discrete spaces, every state is associated with a **neighborhood** that contains other states
- We move through state space, by selecting a neighbor

Hill-Climbing

```
current ← select random initial state
do
  neighbor ← neighbor of current with highest value
  if value(neighbor) ≤ value(current)
    return current
  current ← neighbor
until termination condition is met
```


- Hill-Climbing can be regarded as a discrete state-space version of gradient ascent
- In each step, we select neighbor with highest value

Hill-Climbing Termination Conditions

```
current ← select random initial state
do
  neighbor ← neighbor of current with highest value
  if value(neighbor) ≤ value(current)
    return current
  current ← neighbor
until termination condition is met
```

- ❑ algorithm may end up in non-global local maximum or plateaux
- ❑ termination condition can bound the maximum number of search steps or search time

Example: 8-Queens Problem



- **State-Space:** board configurations with one queen per column
(tuple (8,3,7,4,2,5,1,6) corresponds to board configuration above)
- **Neighborhood:** configurations that can be obtained by moving a single queen to another field in the same column
(hence, every state has $8 * 7 = 56$ neighbors)

Objective Function


- **Objective Function:** number of pairs of queens that can attack each other directly or indirectly (*maximize negative objective function*)

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

- Queen in column 1 **can directly attack** queen in column 2
- Queen in column 1 **can indirectly attack** queen in column 3
- State has **value** -17
- Numbers show **values of neighbors**


Local Optima

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18





- ❑ Hill Climbing goes from left state with value -17 to right state with value -1 (*queen 4 attacks queen 7*) in only 5 steps
- ❑ However, right state is only locally optimal

Excercise: Graph Coloring



- Problem: assign a color from $\{1,2,3\}$ to each node such that no two adjacent nodes have the same color
- Define
 - State space
 - Neighborhood
 - Objective function
- Perform Hill-climbing search from an initial state where all nodes have color 1


Solution: Local Search Problem



- State Space: each assignment of colors from $\{1, 2, 3\}$ to nodes. Formally, we can use 10-tuples like $(1, 1, 1, 1, 1, 1, 1, 1, 1, 1)$
- Neighborhood: change color of one node (each state has $10 * 2 = 20$ neighbors)
- Objective function: number of pairs of adjacent nodes with equal colors (maximize negative objective function)

Solution: Hill-Climbing Search

- Objective function can be increased by at most 3




State (1,1,1,1,1,1,1,1,1,1)

Objective function: -15

Solution: Hill-Climbing Search

- Objective function can be increased by at most 3




State (2,1,1,1,1,1,1,1,1,1)

Objective function: -12

Solution: Hill-Climbing Search

- Objective function can be increased by at most 3




State (2,1,1,1,1,1,1,1,3,1)

Objective function: -9

Solution: Hill-Climbing Search

- Objective function can be increased by at most 2




State (2,1,1,1,2,1,1,1,3,1)

Objective function: -6

Solution: Hill-Climbing Search

- Objective function can be increased by at most 2




State (2,1,1,1,2,3,1,1,3,1)

Objective function: -4

Solution: Hill-Climbing Search

- Objective function can be increased by at most 2




State (2,1,1,1,2,3,1,1,3,2)

Objective function: -2

Solution: Hill-Climbing Search


- Objective function cannot be increased anymore



State (2,1,3,1,2,3,1,1,3,2)

Objective function: 0 (global maximum)

Variants of Hill-Climbing



- Hill-Climbing in its basic form **can perform poorly** in many problems because there is a high risk of ending up in non-global local maxima
- There exist several **variants** that can alleviate the problem

Stochastic Hill-Climbing

```
current ← select random initial state
do
  neighbor ← random neighbor of current with higher value
  if neighbor = null
    return current
  current ← neighbor
until termination condition is met
```

- Instead of selecting neighbor with maximum value, we select random neighbor that improves value
- Probability of selection can increase with value
- Compromise between random search and Hill-Climbing


Other Variants

- **First-choice Hill-climbing:** in neighbor selection step, pick first neighbor that improves objective function
(also useful if neighborhood is too large to enumerate all neighbors)
- **Random-restart (or parallel) Hill-Climbing:** perform n Hill-climbing searches starting from randomly generated initial states
(as n goes to infinity, probability of finding global optimum goes to 1)

A hand holding a game controller is visible in the background, slightly faded. A thick red curved line arches across the top of the slide. The text "Simulated Annealing" is centered in a dark red font.

Simulated Annealing

Downhill Moves



- ❑ Hill-climbing never makes downhill moves
- ❑ However, as figure illustrates, downhill moves are sometimes necessary to find global maximum

Simulated Annealing Intuition

- ❑ **Simulated Annealing** is a randomized algorithm that allows downhill moves
- ❑ In the beginning, the probability for downhill moves is high (exploration)
- ❑ As search progresses, the probability decreases (exploitation)
- ❑ This process is modeled by means of a **temperature variable** that decreases (thus simulated annealing)

Simulated Annealing

```
current ← select random initial state
for  $t = 1$  to  $\infty$ 
     $T \leftarrow \text{schedule}(t)$ 
    if  $T = 0$ 
        return current
    next ← randomly selected neighbor of current
     $\Delta E \leftarrow \text{value}(\textit{next}) - \text{value}(\textit{current})$ 
    if  $\Delta E > 0$ 
        current ← next
    else
        current ← next only with probability  $\exp(\Delta E/T)$ 
```


Cooling Schedule


- `schedule(t)` controls how fast temperature reaches 0
- A slow **cooling schedule** increases probability of finding a high-quality solution but increases runtime
- Some simple **scheduling schemes**:
 - **Stepwise Linear**: start with arbitrary T and decrease T by a constant c in each step
 - **Delayed Stepwise Linear**: start with arbitrary T and decrease T by a constant c every k -th step
- Starting with $T=\infty$ may also make sense to perform random exploration in the beginning

Neighbor Selection

- In each step, Simulated Annealing picks random neighbor
 - If neighbor improves objective, neighbor replaces current state
 - Otherwise, it replaces current state only with probability $\exp(\Delta E/T)$,
where $\Delta E = \text{value}(\text{next}) - \text{value}(\text{current})$

Neighbor Selection: ΔE

- Note that ΔE is always non-positive in else-branch
- Therefore, $0 \leq \exp(\Delta E/T) \leq \exp(0) = 1$
- Probability decreases exponentially with respect to difference
- For instance, for $T=1$, we have
 - $\Delta E = 0 : \exp(0/1) = \exp(0) = 1$
 - $\Delta E = -1 : \exp(-1/1) = 1/e = 0.37$
 - $\Delta E = -2 : \exp(-2/1) = 1/e^2 = 0.14$



Neighbor Selection: Temperature

□ As temperature T decreases, probability further decreases

□ For instance, assume $\Delta E = -2$


□ $T = \infty : \exp(-2/\infty) = \exp(0) = 1$

□ $T=100: \exp(-2/100) = 0.98$

□ $T=10: \exp(-2/10) = 0.81$


□ $T=1: \exp(-2/1) = 0.13$

□ $T=0.1: \exp(-2/0.1) = 0.002$





□ Therefore, Simulated Annealing is similar to **random exploration** for high temperatures and similar to **First-choice Hill-Climbing** later

Exercise: Local Optima Revisited





- As we observed before, given state is locally optimal with value -1
- Compute ΔE for the neighbors obtained by
 - moving queen 2 to row 1 (State (8,1,7,4,2,5,1,6))
 - moving queen 4 to row 8 (State (8,3,7,8,2,5,1,6))
- Compute the probability of performing these moves for $T=10$ and $T=1$

Solution: Local Optima Revisited



- When moving queen 2 to row 1, we add two new attacks
- We have $\Delta E = -3 - (-1) = -2$
- The probabilities of performing these move are
 - $\text{Exp}(-2/10) = 0.81$ for $T=10$
 - $\text{Exp}(-2/1) = 0.13$ for $T=1$

Solution: Local Optima Revisited




- When moving queen 4 to row 8, the value is -2
- We have $\Delta E = -2 - (-1) = -1$
- The probabilities of performing this move are
 - $\text{Exp}(-1/10) = 0.9$ for $T=10$
 - $\text{Exp}(-1/1) = 1/e = 0.37$ for $T=1$



Local Beam Search

Local Beam Search Intuition

- Instead of following a single line through the state space, **Local Beam Search** follows multiple lines (a beam)



- In each step, we select the k best states from the set of all neighbors of all k current states

Local Beam Search

```
k_current ← select k random states
do
  neighbors ← all neighbors of all current states
  k_best_neighbors ← best k states from neighbors
  if no neighbor improves current value
    return k_current
  k_current ← k_best_neighbors
until termination condition is met
```

- Termination conditions can be chosen as before

Local Beam Search vs Parallel Hill-Climbing

- Local Beam Search is **not equivalent** to performing k Hill-Climbing searches in parallel!
- In each step, Local Beam Search **selects the k best neighbors of all current states**
- For instance, if all best neighbors are neighbors of the first current state, we will only use neighbors of this state in the following step

Exercise: 4-Queens Problem

4	♔	♔	4
♔	2	3	4
2	1	2	3
2	1	2	♔

♔	♔	♔	♔
4	5	5	4
4	4	4	4
4	3	3	4

Consider the two states above with the given number of attacks for their neighbors.

- Give two possible next states when performing parallel Hill-climbing (2 runs in parallel).
- Give two possible next states when performing Local Beam Search ($k=2$).

Solution: Hill-Climbing

4	♔	♔	4
♔	2	3	4
2	1	2	3
2	1	2	♔

♔	♔	♔	♔
4	5	5	4
4	4	4	4
4	3	3	4

- Each Hill-climbing search will select a best neighbor for each of the states
- We will continue with one state of value -1 and one state with value -3

Solution: Local Beam Search

4	♔	♔	4
♔	2	3	4
2	1	2	3
2	1	2	♔

♔	♔	♔	♔
4	5	5	4
4	4	4	4
4	3	3	4

- Local Beam Search will select the best two neighbors from the set of all neighbors of both states
- We will continue with two states of value -1

Stochastic Beam Search

- Local Beam Search can concentrate on a small region of the search space too early
- **Stochastic Beam Search** alleviates this problem by using randomized selection similar to Stochastic Hill-Climbing

```
k_current ← select k random states
do
  neighbors ← all neighbors of all current states
  k_best_neighbors ← k ‘random’ states from neighbors
  if no neighbor improves current value
    return k_current
  k_current ← k_best_neighbors
until termination condition is met
```


Random Selection

```
k_current ← select k random states
do
  neighbors ← all neighbors of all current states
  k_best_neighbors ← k 'random' states from neighbors
  if no neighbor improves current value
    return best state from k_current
  k_current ← k_best_neighbors
until termination condition is met
```

- Again, random selection is usually not completely random
- probability of selecting a state should increase with its value
- In this way, we balance diversity (exploration) and exploitation

A hand holding a robotic gripper, symbolizing the intersection of biology and technology. The background is a light blue gradient with a thick red curved line at the top. The text 'Genetic Algorithms' is centered in a dark red serif font.


Genetic Algorithms

Genetic Algorithms Motivation


- Local Beam Search bears some resemblance to **natural selection**:
 - Successors (offspring) of states (organisms) populate next generation
 - Likelihood of survival depends on value (fitness)
- **Genetic Algorithms** extend this analogy
 - In each step, two selected individuals **reproduce**
 - Selection probability increases with fitness (value)
 - There is a small probability of **mutation**
 - Offspring usually replaces other individuals (that die)

Population (States)

- Genetic Algorithms work on a **population** of **individuals**
- Each individual is a state represented as a **chromosome** (e.g. string) over a set of **genes** (e.g. set of characters)
- For 8-Queens problem, we could use string consisting of 8 integers from $\{1, \dots, 8\}$, where i -th integer is row position of i -th queen




83742516




31645372

Fitness (Objective Function)

- **Fitness** corresponds to objective function
- For 8-Queens problem, we can reuse our old value function
- In order to get a non-negative fitness function, we could count the number of pairs of queens that cannot attack each other
(28 – number of attacks between pairs)




Fitness -1 or 27 (28-1)



Fitness -6 or 22 (28-6)


Reproduction (Recombination)

- **Reproduction:** we choose a random **crossover point** in chromosome
- Offspring is created by crossing parents at this point




Mutation

- With small probability, **mutation** of offspring occurs
- E.g. replace one gene in chromosome randomly



837 45372



837 41372

A Simple Genetic Algorithm

repeat

population \leftarrow create k chromosomes at random

for i = 1 **to** k **do**

x \leftarrow select random chromosome based on fitness

y \leftarrow select random chromosome based on fitness

child \leftarrow reproduce(x , y)

if (random() < 0.05)

 child \leftarrow mutate(child)


 add child to population

 remove random chromosome based on fitness

until some chromosome is fit enough, or time limit is reached


return the best chromosome in population

Excercise: Graph Coloring




- Problem: assign a color from $\{1,2,3\}$ to each node such that no two adjacent nodes have the same color
- Define
 - Genes
 - Chromosomes
 - Fitness (value) of individuals
 - Mutation operation
- Illustrate reproduction and mutation by means of a small example

Solution: Representation




- Genes {R,G,B}
- Chromosomes are strings of length 10, where i-th gene (letter) represents color of i-th node
- Fitness is the number of edges such that the corresponding nodes have different colors ($0 \leq \text{Fitness} \leq 15$)
- Mutation operation changes color of random node (50% chance for each of the other two colors)

Solution: Reproduction



RBBRGG RGRB

RBBRGG GRBG




RGBGRB GRBG

RGBGRB RGRB


Solution: Mutation

- Mutation at 9-th gene (node 8)

RBBRGG GRBG



RBBRGG GRRG



Set Cover Problem

- **Set Cover Problem:** given set of n elements $U = \{e_1, \dots, e_n\}$ and subsets S_1, \dots, S_m of U , find a minimal number of subsets that contain all elements from U (a **minimal set cover** of U)
- For example, U can be a shopping list and subsets correspond to shops that offer the desired items
 - $U = \{1, 2, 3, 4, 5\}$
 - $S_1 = \{1, 2, 5\}$, $S_2 = \{1, 4\}$, $S_3 = \{3, 5\}$, $S_4 = \{1, 2\}$, $S_5 = \{3, 4\}$
 - S_2 , S_3 and S_4 contain all elements from U (S_2 , S_3 and S_4 form a set cover of size 3)
 - S_1 and S_5 do also contain all elements (S_1 and S_5 form a set cover of size 2)
 - Hence, the set cover S_2 , S_3 and S_4 is not minimal

Excercise: Set Cover Problem

- **Set Cover Problem:** given set of n elements $U = \{e_1, \dots, e_n\}$ and m subsets S_1, \dots, S_m of U , find a minimal number of subsets that contain all elements from U
- Define
 - Genes
 - Chromosomes
 - Fitness (value) of individuals
 - Mutation operation
- Illustrate reproduction by means of a small example

Solution: Representation

- Genes $\{0, 1\}$
- Chromosomes are strings of length m , where i -th gene (letter) indicates that i -th set is used
- $U = \{1, 2, 3, 4, 5, 6\}$
- $S_1 = \{1, 2, 5\}$, $S_2 = \{1, 4, 6\}$, $S_3 = \{2, 3, 5\}$, $S_4 = \{1, 2, 6\}$, $S_5 = \{3, 4\}$
- 10101: use S_1 , S_3 and S_5
- 01101: use S_2 , S_3 and S_5

Solution: Fitness

□ Fitness of an individual is

□ -1 if individual does not represent a set cover

□ m - size of set cover (number of 1s) otherwise

□ $U = \{1, 2, 3, 4, 5, 6\}$

□ $S_1 = \{1, 2, 5\}$, $S_2 = \{1, 4, 6\}$, $S_3 = \{2, 3, 5\}$, $S_4 = \{1, 2, 6\}$, $S_5 = \{3, 4\}$

□ 10101 covers $\{1, 2, 3, 4, 5\}$: fitness -1 (does not cover 6)

□ 01101 covers $\{1, 2, 3, 4, 5, 6\}$: fitness $5 - 3 = 2$


□ 01100 covers $\{1, 2, 3, 4, 5, 6\}$: fitness $5 - 2 = 3$

Solution: Reproduction and Mutation

- Mutation operation flips random gene



Excercise: Traveling Salesman Problem




	0	1	2	3	4
0	0	5	6	5	2
1	3	0	6	2	1
2	4	3	0	5	1
3	5	5	3	0	7
4	6	6	4	7	0

(i-th column shows cost from node i to other nodes)


- ☐ **TSP**: find cyclic route that starts from node 0, visits each node exactly once and minimizes the overall edge cost
- ☐ Define
 - ☐ Genes
 - ☐ Chromosomes
 - ☐ Fitness (value) of individuals
 - ☐ Mutation operation
- ☐ Does crossover operation make sense?

Solution: Representation

- Genes {1, 2, 3, 4}
- Chromosomes are strings of length 4, where i -th gene (letter) represents the node that is visited at time i (first node is fixed to be 0)



1234




3124

Solution: Fitness

	0	1	2	3	4
0	0	5	6	5	2
1	3	0	6	2	1
2	4	3	0	5	1
3	5	5	3	0	7
4	6	6	4	7	0

- Fitness is 35 (rough upper bound on max. cost) minus the cost of getting from i -th to $(i+1)$ -th node and from fifth node back to first node





1234


- Cost from 0 to 1: 3
- Cost from 1 to 2: 3
- Cost from 2 to 3: 3
- Cost from 3 to 4: 7
- Cost from 4 to 0: 2
- Overall cost: 18
- **Fitness:** $35 - 18 = 17$

Solution: Mutation

- Mutation operation switches two genes at random
- E.g. switch first and third gene



Solution: Problem with Reproduction



Naive reproduction yields invalid solutions

Reproduction of Permutations

- ❑ If chromosomes correspond to **permutations** like in TSP, naive crossover reproduction can yield invalid solutions
- ❑ the problem can be fixed by
 - ❑ Changing the **representation**
 - ❑ Designing an additional **repair operation** that is applied after recombination
 - ❑ Using special **crossover techniques for permutations**


Uniform-Order Crossover

□ Uniform-Order Crossover

- Select two parents at random
- Create **random binary template**
- Child 1 is created by using genes of parent 1 at **1-positions**. Fill gaps with the remaining elements according to the order given by parent 2
- Child 2 is created by switching the roles of parent 1 and parent 2

A	B	C	D	E	F	G	Parent P_1
E	B	D	C	F	G	A	Parent P_2
0	1	1	0	0	1	0	Template
E	B	C	D	G	F	A	Child C_1
A	B	D	C	E	G	F	Child C_2

Solution: Problem with Reproduction



Variants

- Genetic Algorithms can be configured by different
 - Selection operators
 - Reproduction operators
 - Mutation operators
- Hybrid Genetic Algorithms combine genetic algorithms with other techniques
 - for instance, Memetic Algorithms apply a (fast) local search algorithm to each newly created individual to make it locally optimal
 - Hill-Climbing is well suited for this purpose because it is fast

A hand holding a game controller is visible in the background, slightly faded. A thick red curved line arches across the top of the slide. The text "Summary: Local Search" is centered in a dark red font.

Summary: Local Search

Summary

- If we are only interested in a solution and not in the solution path, **Local Search** can be better suited than Classical Search
- this is typically the case for **optimization** and **constraint satisfaction** problems
- Some **important local search algorithms** are
 - Hill-Climbing
 - Simulated Annealing
 - Local Beam Search
 - Genetic Algorithms

Further Readings

Lecture is mainly based on:

Russell, S., Norvig, P. Artificial Intelligence - A modern approach. Pearson Education: 2010.

More details on presented (and similar algorithms) can be found in:

Burke, E. K., & Kendall, G. Search methodologies - Introductory Tutorials in Optimization and Decision Support Techniques. Springer US: 2014.